


Matrix Factorization

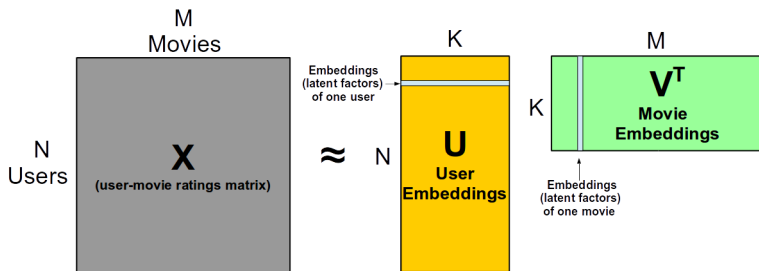
- Given a matrix \mathbf{X} of size $N \times M$, approximate it as a product of two matrices



- \mathbf{U} : $N \times K$ latent factor matrix
 - Each row of \mathbf{U} represents a K -dim latent factor \mathbf{u}_n
- \mathbf{V} : $M \times K$ latent factor matrix
 - Each row of \mathbf{V} represents a K -dim latent factor \mathbf{v}_n
- Each entry of \mathbf{X} can be written as: $X_{nm} \approx \mathbf{u}_n^T \mathbf{v}_m = \sum_{k=1}^K u_{nk} v_{mk}$
- If X_{nm} is large (small) then \mathbf{u}_n and \mathbf{v}_m should be similar (dissimilar)

Why Matrix Factorization?


- The latent factors can be used/interpreted as “embeddings” or “features”



- Especially useful for learning good features for “dyadic” or relational data
 - Examples: Users-Movies ratings, Users-Products purchases, etc.
- If $K \ll \min\{M, N\} \Rightarrow$ then can also be seen as dimensionality reduction or a “low-rank factorization” of the matrix \mathbf{X}

Why Matrix Factorization?

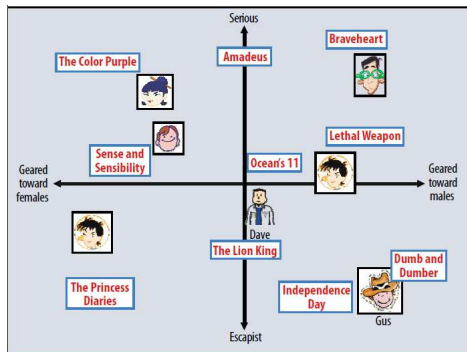
- Can also predict the missing/unknown entries in the original matrix



- Note: The latent factor matrices \mathbf{U} and \mathbf{V} can be learned even when the matrix \mathbf{X} is only **partially observed** (as we will see shortly)
- After learning \mathbf{U} and \mathbf{V} , any missing X_{nm} can be approximated by $\mathbf{u}_n^T \mathbf{v}_m$
- \mathbf{UV}^T is the best low-rank matrix that approximates the full \mathbf{X}
- Note: The “**Netflix Challenge**” was won by a matrix factorization method

Interpreting the Embeddings/Latent Factors


- Embeddings/latent factors can often be interpreted. E.g., as “genres” if \mathbf{X} represents a user-movie rating matrix. A cartoon with $K = 2$ shown below



- Similar things (users/movies) get embedded nearby in the embedding space (two things will be deemed similar if their embeddings are similar). Thus useful for **computing similarities** and/or **making recommendations**


Matrix Factorization

- Recall our matrix factorization model: $\mathbf{X} \approx \mathbf{U}\mathbf{V}^T$
- Goal: learn \mathbf{U} and \mathbf{V} , given a subset Ω of entries in \mathbf{X} (let's call it \mathbf{X}_Ω)
- Some notations:
 - $\Omega = \{(n, m)\}$: X_{nm} is observed
 - Ω_{r_n} : column indices of observed entries in row n of \mathbf{X}
 - Ω_{c_m} : row indices of observed entries in column m of \mathbf{X}



Matrix Factorization

- We want \mathbf{X} to be as close to \mathbf{UV}^T as possible



- Let's define a squared “loss function” over the observed entries in \mathbf{X}

$$\mathcal{L} = \sum_{(n,m) \in \Omega} (x_{nm} - \mathbf{u}_n^T \mathbf{v}_m)^2$$

- Here the latent factors $\{\mathbf{u}_n\}_{n=1}^N$ and $\{\mathbf{v}_m\}_{m=1}^M$ are the **unknown parameters**
- Squared loss chosen only for simplicity; other loss functions can be used
- How do we learn $\{\mathbf{u}_n\}_{n=1}^N$ and $\{\mathbf{v}_m\}_{m=1}^M$?

Alternating Optimization

- We will use an ℓ_2 regularized version of the squared loss function

$$\mathcal{L} = \sum_{(n,m) \in \Omega} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \sum_{n=1}^N \lambda_U \|\mathbf{u}_n\|^2 + \sum_{m=1}^M \lambda_V \|\mathbf{v}_m\|^2$$

- A **non-convex** problem. Difficult to optimize w.r.t. \mathbf{u}_n and \mathbf{v}_m jointly.
- One way is to solve for \mathbf{u}_n and \mathbf{v}_m in an **alternating fashion**, e.g.,

- $\forall n$, fix all variables except \mathbf{u}_n and solve the optim. problem w.r.t. \mathbf{u}_n

$$\arg \min_{\mathbf{u}_n} \sum_{m \in \Omega_{r_n}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \|\mathbf{u}_n\|^2$$

- $\forall m$, fix all variables except \mathbf{v}_m and solve the optim. problem w.r.t. \mathbf{v}_m

$$\arg \min_{\mathbf{v}_m} \sum_{n \in \Omega_{c_m}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_V \|\mathbf{v}_m\|^2$$

- Iterate until not converged
- Each of these subproblems has a simple, convex objective function
- Convergence properties of such methods have been studied extensively

The Solutions

- Easy to show that the problem

$$\arg \min_{\mathbf{u}_n} \sum_{m \in \Omega_{r_n}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \|\mathbf{u}_n\|^2$$

.. has the solution

$$\mathbf{u}_n = \left(\sum_{m \in \Omega_{r_n}} \mathbf{v}_m \mathbf{v}_m^\top + \lambda_U \mathbf{I}_K \right)^{-1} \left(\sum_{m \in \Omega_{r_n}} X_{nm} \mathbf{v}_m \right)$$

- Likewise, the problem

$$\arg \min_{\mathbf{v}_m} \sum_{n \in \Omega_{c_m}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_V \|\mathbf{v}_m\|^2$$

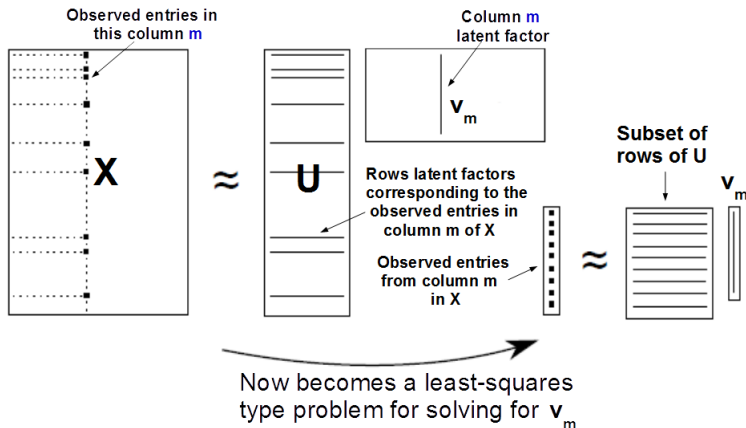
.. has the solution

$$\mathbf{v}_m = \left(\sum_{n \in \Omega_{c_m}} \mathbf{u}_n \mathbf{u}_n^\top + \lambda_V \mathbf{I}_K \right)^{-1} \left(\sum_{n \in \Omega_{c_m}} X_{nm} \mathbf{u}_n \right)$$

- Note that this is very similar to (regularized) least squares regression
- Thus matrix factorization can be also seen as [a sequence of regression problems](#) (one for each latent factor)

Matrix Factorization as Regression

Suppose we are solving for \mathbf{v}_m (with \mathbf{U} and all other \mathbf{v}_m 's fixed)



Can think of solving for \mathbf{u}_n (with \mathbf{V} and all other \mathbf{u}_n 's fixed) in the same way

Matrix Factorization as Regression

- A very useful way to understand matrix factorization
- Can modify the regularized least-squares like objective

$$\arg \min_{\mathbf{u}_n} \sum_{m \in \Omega_{r_n}} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \mathbf{u}_n^\top \mathbf{u}_n$$

.. using other **loss functions** and **regularizers**

- Some possible modifications:
 - If entries in the matrix \mathbf{X} are binary, counts, etc. then we can replace the squared loss function by some other loss function (e.g., logistic or Poisson)
 - Can impose other constraints on the latent factors, e.g., non-negativity, sparsity, etc. (by changing the regularizer)
 - Can think of this also as a probabilistic model (a likelihood function on X_{nm} and priors on the latent factors $\mathbf{u}_n, \mathbf{v}_m$) and do MLE/MAP

Matrix Factorization: The Complete Algorithm

- Input: Partially complete matrix \mathbf{X}_Ω
- Initialize the latent factors $\mathbf{v}_1, \dots, \mathbf{v}_M$ randomly
- Iterate until converge
 - Update each row latent factor \mathbf{u}_n , $n = 1, \dots, N$ (can be in parallel)

$$\mathbf{u}_n = \left(\sum_{m \in \Omega_{r_n}} \mathbf{v}_m \mathbf{v}_m^\top + \lambda_U \mathbf{I}_K \right)^{-1} \left(\sum_{m \in \Omega_{r_n}} X_{nm} \mathbf{v}_m \right)$$

matmul(X,V)

- Update each column latent factor \mathbf{v}_m , $m = 1, \dots, M$ (can be in parallel)

$$\mathbf{v}_m = \left(\sum_{n \in \Omega_{c_m}} \mathbf{u}_n \mathbf{u}_n^\top + \lambda_V \mathbf{I}_K \right)^{-1} \left(\sum_{n \in \Omega_{c_m}} X_{nm} \mathbf{u}_n \right)$$

- Final prediction for any entry: $X_{nm} = \mathbf{u}_n^\top \mathbf{v}_m$

A Faster Algorithm via SGD

- Alternating optimization is nice but can be slow (note that it requires matrix inversion with cost $O(K^3)$ for updating each latent factor $\mathbf{u}_n, \mathbf{v}_m$)
- An alternative is to use stochastic gradient descent (SGD). In each round, select a randomly chosen entry X_{nm} with $(n, m) \in \Omega$
- Consider updating \mathbf{u}_n . For loss function $\sum_{m \in \Omega_n} (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m)^2 + \lambda_U \|\mathbf{u}_n\|^2$, the **stochastic gradient** w.r.t. \mathbf{u}_n using this randomly chosen entry X_{nm} is

$$-(X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m) \mathbf{v}_m + \lambda_U \mathbf{u}_n$$

- Thus the SGD update for \mathbf{u}_n will be

$$\mathbf{u}_n = \mathbf{u}_n - \eta(\lambda_U \mathbf{u}_n - (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m) \mathbf{v}_m)$$

- Likewise, the SGD update for \mathbf{v}_m will be

$$\mathbf{v}_m = \mathbf{v}_m - \eta(\lambda_V \mathbf{v}_m - (X_{nm} - \mathbf{u}_n^\top \mathbf{v}_m) \mathbf{u}_n)$$

- The SGD algorithm chooses a random entry X_{nm} in each iteration, updates $\mathbf{u}_n, \mathbf{v}_m$, and repeats until convergece (\mathbf{u}_n 's, \mathbf{v}_m 's randomly initialized).